



How the system runs — sources, Python, LLM, and where you click

Every job is **manually triggered** (a CLI command or a dashboard button) — nothing is auto-scheduled. Inside a job, Python steps run automatically; LLM steps run inside a Claude session; capital decisions stop at a human click.



MANUAL vs AUTOMATIC — nothing runs on a schedule. You start each job by hand:

- **Daily (no tokens):** `--job daily_data` → all ☞ Python pipelines → DB. Or click **Refresh data** on the dashboard.
- **Monthly (no tokens):** `--job monthly_macro / monthly_amc_portfolios / monthly_newsletter_digest`.
- **Screening (staged):** `--job universe_screening_stage` (~100 names/run, ~7-8 runs).
- **Ideas & research (LLM, in a Claude session):** `--job weekly_cycle_assessment` then `weekly_idea_cycle`; brain prepares a packet per step, you run the agent, output is validated and re-ingested.
- **You click to decide:** `--decision APPROVE|REJECT|MODIFY` — the only thing that changes the portfolio.
- **Manual downloads (dashed red):** VAHAN vehicle regs, EPFO/NPCI, and HDFC/SBI/Kotak AMC files are bot/JS-walled — one Playwright install unlocks them; until then, hand-download + `macro_manual` import.